

# **Red-Black Trees**

Data Structures and Algorithms  
Andrei Bulatov

## Red-Black Trees

All binary search tree operations take  $O(h)$  time, where  $h$  is the height of the tree

Therefore, it is important to 'balance' the tree so that its height is as small as possible

There are many ways to achieve this

One of them: **Red-Black trees**

Every node of such a tree contains one extra bit, its color

Another agreement: the Nil pointer is treated as a leaf, an extra node


The rest of the nodes are called **internal**

## Red-Black Properties


A binary search tree is a red-black tree if it satisfies the following red-black properties:

- Every node is either red or black
- The root is black
- Every leaf (Nil) is black
- If a node is red, then both its children are black
- For each node, all paths from the node to descendant leaves contain the same number of black nodes

## Example



## Example



## Black Height

The number of black nodes on paths from node  $x$  to its descendant leaves in a red-black tree is called its **black height**, denoted  $bh(T)$

### Lemma

A red-black tree with  $n$  internal nodes has height at most  $2 \cdot \log(n + 1)$

### Proof

We show first that the subtree rooted at  $x$  contains at least  $2^{bh(x)} - 1$  nodes

Induction on  $bh(x)$

Base Case: If  $bh(x) = 0$ , then  $x$  is a leaf,  $nil[T]$

In this case,  $2^{bh(x)} - 1 = 2^0 - 1 = 0$  internal nodes

## Black Height (cntd)

Inductive hypothesis: the claim is true for any  $y$  with height less than that of  $x$

Inductive Case: Let  $bh(x) > 0$  and  $x$  has two children

The black height of the children is either  $bh(x)$  or  $bh(x) - 1$ , depending on its color

Since the children have smaller height, we can apply the induction hypothesis

Thus the subtree rooted at  $x$  contains at least

$$(2^{bh(x)-1} - 1) + (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$$

nodes

## Black Height (cntd)

Suppose  $h$  is the height of the tree

By the red-black property at least half of nodes on every root-to-leaf path are black (not including the root)

Therefore the black height of the root is at least  $h/2$

Thus

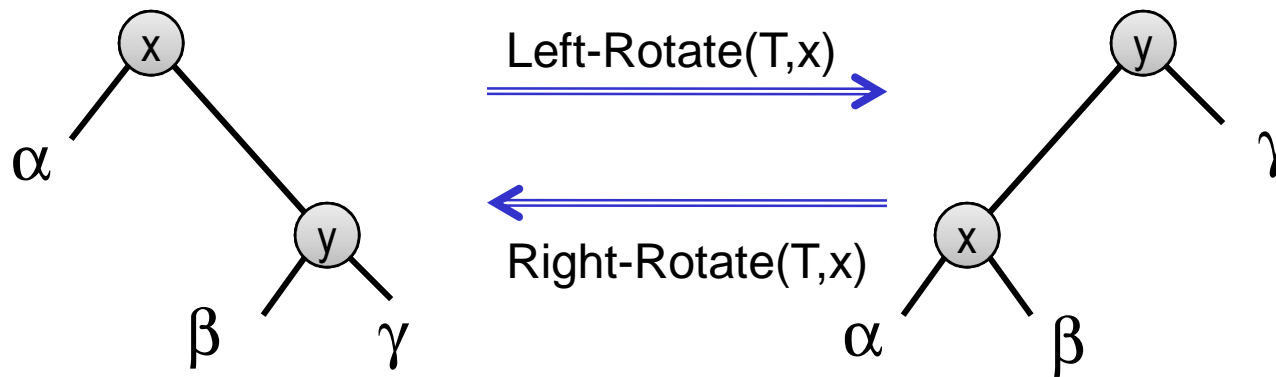
$$n \geq 2^{h/2} - 1$$
$$\log(n+1) \geq h/2$$

QED



## Rotations


Sometimes we will need to rearrange pointers inside an RB-tree



## Rotations: Pseudocode

```
Left-Rotate(T,x)
  set y:=right[x]
  set right[x]:=left[y]
  set parent[left[y]]:=x
  set parent[y]:=parent[x]
  if parent[x]=nil[T] then
    set root[T]:=y
  else if x=left[parent[x]] then
    set left[parent[x]]:=y
  else
    right[parent[x]]:=y
  set left[y]:=x
  set parent[x]:=y
```

## Rotations: Example



## Insertion

Insertion for RB-trees is done in the same way as for ordinary binary search trees.

Except:

- we should be careful about Nil links

- the new node is colored red

- the resulting tree may not be an RB-tree, we need to fix it

## Insertion: Pseudocode

RB-Insert( $T, z$ )

set  $y := \text{Nil}[T]$ ,  $x := \text{root}[T]$

while  $x \neq \text{Nil}[T]$  do

    set  $y := x$

    if  $\text{key}[z] < \text{key}[x]$  then set  $x := \text{left}[x]$

        else set  $x := \text{right}[x]$

endwhile

set  $\text{parent}[z] := y$

if  $y = \text{Nil}[T]$  then set  $\text{root}[T] := z$


    else if  $\text{key}[z] < \text{key}[y]$  then set  $\text{left}[y] := z$

        else set  $\text{right}[y] := z$

set  $\text{left}[z] := \text{Nil}[T]$   $\text{right}[z] := \text{Nil}[T]$   $\text{color}[z] := \text{RED}$

RB-Insert-FixUp( $T, z$ )

# Insertion: FixUp



## FixUp: Pseudocode

RB-Insert-FixUp(T,z)

```
while color[parent[z]]=RED do
  if parent[z]=left[parent[parent[z]]] then do
    set y:=right[parent[parent[z]]]
    if color[y]=RED then do
      set color[parent[z]]:=BLACK  color[y]:=BLACK
      set color[parent[parent[z]]]:=RED
      set z:=parent[parent[z]]
    else do
      if z=right[parent[z]] then do
        set z:=parent[z]    Left-Rotate(T,z)
      set color[parent[z]]:=BLACK
      set color[parent[parent[z]]]:=RED
    else  (same as then with left and right swopped)
color[root[T]]:=BLACK
```

# Soundness

## Lemma

Given an RB-tree, algorithm RB-Insert produces an RB-tree

## Proof

We show that the main loop preserves the following loop invariant:

- (a) Node  $z$  is red
- (b) If  $\text{parent}[z]$  is the root then  $\text{parent}[z]$  is black
- (c) If there is a violation of the red-black property, there is at most one violation, and it is of either property 2 or property 4.

If there is a violation of property 2, it occurs because  $z$  is the root and red

If there is a violation of property 4, it occurs because both  $z$  and  $\text{parent}[z]$  are red



## Soundness: Initialization

- (a) When `RB-Insert-FixUp` is called,  $z$  is a red node
- (b) If  $p[z]$  is the root, it hasn't changed yet, and so is black
- (c) Properties 1, 3, and 5 are true

If property 2 is violated, then the red root is just added, that is the root is  $z$ . In this case the tree does not have internal vertices.

Property 4 is not violated in this case

If property 4 is violated, then, since the children of  $z$  are black sentinels and the tree had no prior violations, the violation must be because both  $z$  and  $\text{parent}[z]$  are red.

There are no other violations of red-black properties.

## Soundness: Termination

When the big loop of `RB-Insert-FixUp` terminates `parent[z]` is black  
(if `z` is the root then `parent[z]` is the sentinel)

Therefore there is no violation of property 4.

If property 2 is violated, then the last line of `RB-Insert-FixUp` restores it.

## Soundness: Maintenance


There are several cases to consider. We only consider those in which  $\text{parent}[z]$  is the left child of  $\text{parent}[\text{parent}[z]]$

This is the situation given in the pseudocode


If  $\text{parent}[\text{parent}[z]]$  does not exist then  $\text{parent}[z]$  is the root, and by part (b) is black.

Therefore if the loop does not terminate here,  $\text{parent}[\text{parent}[z]]$  exists

**Case 1.**  $z$ 's uncle  $y$  is red



OR




## Soundness: Maintenance (cntd)

Both  $y$  and  $\text{parent}[z]$  are red, and  $\text{parent}[\text{parent}[z]]$  is black

We color  $y$  and  $\text{parent}[z]$  black, and  $\text{parent}[\text{parent}[z]]$  – red

We then set  $\text{parent}[\text{parent}[z]]$  to be the new node  $z$



Show that the loop invariant is preserved

- (a) Since we color  $\text{parent}[\text{parent}[z]]$  red, the new  $z$  (denoted  $z'$ ) is red

## Soundness: Maintenance (cntd)

- (b) We have  $\text{parent}[z'] = \text{parent}[\text{parent}[\text{parent}[z]]]$ , and the color of this node does not change.

If it is the root, it is black

- (c) It is easy to see that Case 1 maintains property 5, and does not violate properties 1 and 3

If  $z'$  is the root, then Case 1 corrected the only violation of property 4. Since  $z'$  is red and it is the root, property 2 becomes violated, and it is due to  $z'$

If  $z'$  is not the root, then Case 1 has not created a violation of property 2.

Case 1 corrected the violation of property 4 due to  $z$  and  $\text{parent}[z]$ .


If  $\text{parent}[z']$  is black then we are done

Otherwise  $z'$  and  $\text{parent}[z']$  create the only violation of property 4

## Soundness: Maintenance (cntd)

**Case 2.** z's uncle y is black and z is a right child

**Case 3.** z's uncle y is black and z is a left child



## Soundness: Maintenance (cntd)

**Case 2.** z's uncle y is black and z is a right child

**Case 3.** z's uncle y is black and z is a left child

In Case 2 the algorithm performs rotation and we get Case 3

Since both nodes being rotated are red no RB properties change

Moreover,  $\text{parent}[\text{parent}[z]]$  exists and it does not change when

Case 2 is transformed into Case 3

In Case 3 we perform same rotations and color changes that does not affect property 5

Then since there are no longer 2 red nodes in a row, the while loop is not executed again, for  $\text{parent}[z]$  is now black

## Soundness: Maintenance (cntd)

- (a) Case 2 makes  $z$  points to  $\text{parent}[z]$ , which is red  
Nothing changes in Case 3
- (b) Case 2 does not change the color of the root  
Case 3 makes  $\text{parent}[z]$  black, so if it is the root, it is black
- (c) As in Case 1, properties 1, 3, and 5 are maintained in Cases 2 and 3

Since node  $z$  is not the root, there is no violation of property 2

Cases 2 and 3 do not introduce any violation of property 2, since the only node that becomes red also becomes a child of a black node by the rotation in Case 3

Cases 2 and 3 correct the only violation of property 4 and do not introduce another violation



## Running Time

Since the height of an RB-tree with  $n$  nodes is  $O(\log n)$ , RB-Insert takes  $O(\log n)$  time, except for RB-Insert-FixUp

In RB-Insert-FixUp the while loop repeats only if Case 1 takes place, and in this case  $z$  moves up 2 levels.

The total number of iterations is  $O(\log n)$

## RB-Deletion


RB-Delete also first runs the regular Deletion algorithm, and then fixes violations of the RB properties

RB-Delete( $T, z$ )



```
if left[z]=Nil[T] or right[z]=Nil[T] then set y:=z
    else set y:=Tree-Successor(z)
if left[y]≠Nil[T] then set x:=left[y]
    else set x:=right[y]
if x≠Nil[T] then set parent[x]:=parent[y]
    else set x:=right[x]
if parent[y]=Nil[T] then set root[T]:=x
    else if y=left[parent[y]]:=x then left[parent[y]]:=x
        else right[parent[y]]:=x

if y≠z then do
    set key[z]:=key[y]
    copy y's data into z
if color[y]=BLACK then  RB-Delete-FixUp(T,x)
return y
```

## Deletion: Pictures



## Deletion: Pictures



new  $x = \text{root}[T]$

## RB-Delete-FixUp

RB-Delete-FixUp(T,x)

while  $x \neq \text{root}[T]$  and  $\text{color}[x] = \text{BLACK}$  do

  if  $x = \text{left}[\text{parent}[x]]$  then do

    set  $w := \text{right}[\text{parent}[x]]$

    if  $\text{color}[w] = \text{RED}$  then do

      set  $\text{color}[w] := \text{BLACK}$      $\text{color}[\text{parent}[x]] := \text{RED}$

      Left-Rotate(T, parent[x])    set  $w := \text{right}[\text{parent}[x]]$

    if  $\text{color}[\text{left}[w]] = \text{BLACK}$  and  $\text{color}[\text{right}[w]] = \text{BLACK}$  then

      set  $\text{color}[w] := \text{RED}$      $x := \text{parent}[x]$

    else if  $\text{color}[\text{right}[w]] = \text{BLACK}$  then do

      set  $\text{color}[\text{left}[w]] := \text{BLACK}$      $\text{color}[w] := \text{RED}$

      Right-Rotate(T, w)    set  $w := \text{right}[\text{parent}[x]]$

      set  $\text{color}[w] := \text{color}[\text{parent}[x]]$

      set  $\text{color}[\text{parent}[x]] := \text{BLACK}$

      set  $\text{color}[\text{right}[w]] := \text{BLACK}$

      Left-Rotate(T, parent[x])    set  $x := \text{root}[T]$

  else (same as then flipping left and right)

set  $\text{color}[x] := \text{BLACK}$

Case 1  
Case 2  
Case 3  
Case 4